

14. SQL queries

Principles of Data Science with R

Dr. Uma Ravat

PSTAT 10

We saw...

- Databases store massive amounts of data that cannot fit in memory.
- SQL(Structured Query Language) is used to manipulate relational databases
- SQLite is the SQL implementation we will use, provided by the RSQLite package.
- Three parts of the relational data model
 1. Manipulative
 - SQL for create, update, delete tables, databases and user access
 - SQL for select, insert, update, delete data in tables
 2. Structural
 - ER diagram, Database schema
 - Primary keys, Foreign keys
 3. Integrity
 - Entity integrity: integrity of each relation
 - Referential integrity : integrity between relations

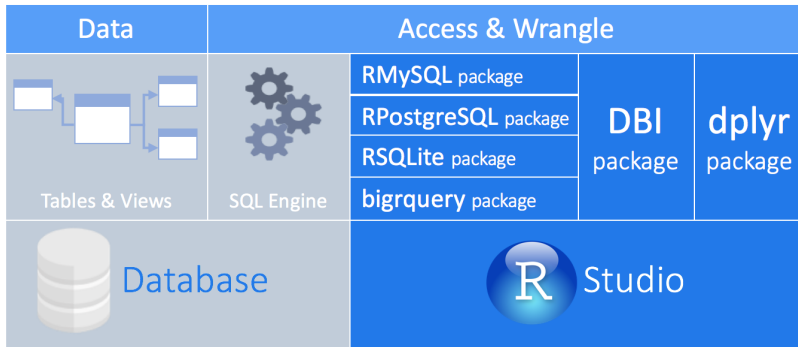
Next we will see. . .

Manipulative Part of relational databases

- More SQL queries
- All parts of a SQL query

SQLite RDBMS and Rstudio

Open Source Databases



ChinookDB Entity-Relationship Diagram (ER Diagram)



Connecting to a DB

```
library(DBI)
library(RSQLite)
drv = dbDriver("SQLite") # the driver for the db you want to connect to
chinook_db = dbConnect(drv, # the driver to use
                        dbname="./data/Chinook_Sqlite.sqlite") #path to the db file
```

chinook_db is an R object that represents a connection to the database file Chinook_Sqlite.sqlite

SQL queries

```
SELECT columns  
  FROM table  
  WHERE condition  
  GROUP BY columns  
  HAVING condition  
  ORDER BY column [ASC | DESC]  
  LIMIT offset, count;
```

WHERE, GROUP BY, HAVING, ORDER BY, LIMIT are all optional

We saw SELECT, FROM, ORDER BY, LIMIT

More SQL queries

Recall SELECT

```
dbGetQuery(chinook_db,  
            "SELECT count(*) FROM track")
```

```
##      count(*)  
## 1          3503
```

What are all the fields for every track?

```
dbListFields(chinook_db, "track")
```

```
## [1] "TrackId" "Name" "AlbumId" "MediaTypeId" "GenreId"  
## [6] "Composer" "Milliseconds" "Bytes" "UnitPrice"
```

```
track_sel <- dbGetQuery(chinook_db,  
                        "SELECT * FROM track")
```

```
str(track_sel)
```

```
## 'data.frame': 3503 obs. of 9 variables:  
## $ TrackId : int 1 2 3 4 5 6 7 8 9 10 ...  
## $ Name : chr "For Those About To Rock (We Salute You)" "Balls to the Wall"  
"Fast As a Shark" "Restless and Wild" ...  
## $ AlbumId : int 1 2 3 3 3 1 1 1 1 1 ...  
## $ MediaTypeId : int 1 2 2 2 2 1 1 1 1 1 ...  
## $ GenreId : int 1 1 1 1 1 1 1 1 1 1 ...  
## $ Composer : chr "Angus Young, Malcolm Young, Brian Johnson" NA "F. Baltes,  
S. Kaufman, U. Dirksneider & W. Hoffman" "F. Baltes, R.A. Smith-Diesel, S.  
Kaufman, U. Dirksneider & W. Hoffman" ...  
## $ Milliseconds: int 343719 342562 230619 252051 375418 205662 233926 210834  
203102 263497 ...  
## $ Bytes : int 11170334 5510424 3990994 4331779 6290521 6713451 7636561  
6852860 6599424 8611245 ...
```

Suppose we only want the first five records for TrackId, Name, AlbumId, Milliseconds, Bytes, UnitPrice from Track table

```
dbGetQuery(chinook_db,  
           "SELECT TrackId, Name, AlbumId, Milliseconds, Bytes, UnitPrice  
           FROM track  
           limit 5")
```

```
## TrackId Name AlbumId Milliseconds Bytes  
## 1 1 For Those About To Rock (We Salute You) 1 343719 11170334  
## 2 2 Balls to the Wall 2 342562 5510424  
## 3 3 Fast As a Shark 3 230619 3990994  
## 4 4 Restless and Wild 3 252051 4331779  
## 5 5 Princess of the Dawn 3 375418 6290521  
## UnitPrice  
## 1 0.99  
## 2 0.99  
## 3 0.99  
## 4 0.99  
## 5 0.99
```

SELECT, expanded

In the first line of SELECT, we can directly specify **computations** that we want performed

```
SELECT columns or computations  
FROM table  
WHERE condition  
GROUP BY columns  
HAVING condition  
ORDER BY column [ASC | DESC]  
LIMIT offset, count;
```

Main tools for computations:

MIN, MAX, COUNT, SUM, AVG or any math formula

Example

To calculate the average Milliseconds, Bytes and Max UnitPrice

```
dbGetQuery(chinook_db,  
  "SELECT AVG(Milliseconds), AVG(Bytes), MAX(UnitPrice)  
  FROM Track")
```

```
##      AVG(Milliseconds) AVG(Bytes) MAX(UnitPrice)  
## 1          393599.2    33510207          1.99
```

To replicate this simple command on an imported data frame:

```
mean(track_sel$Milliseconds, na.rm=TRUE)
```

```
## [1] 393599.2
```

```
mean(track_sel$Bytes, na.rm=TRUE)
```

```
## [1] 33510207
```

```
max(track_sel$UnitPrice, na.rm=TRUE)
```

```
## [1] 1.99
```

GROUP BY

We can use the GROUP BY option in SELECT to define aggregation groups

The GROUP BY statement groups rows that have the same values and performs aggregation (COUNT(), MAX(), MIN(), SUM(), AVG()) on those groups.

The diagram illustrates the process of grouping data by AlbumId and calculating the average Bytes for each group.

Table 1: "select AlbumId, Bytes from track"

AlbumId	Bytes
2	5510424
3	3990994
1	6713451
1	7636561
4	10847611
5	9719579
5	10552051
1	6852860
1	6599424
3	4331779
3	6290521
4	7032162
4	12021261
4	8776140
5	8675345
5	10144730

Table 2: groupby AlbumId

AlbumId	Bytes
1	6713451
1	7636561
1	6852860
1	6599424
2	5510424
3	3990994
3	4331779
3	6290521
4	10847611
4	7032162
4	12021261
4	8776140
5	9719579
5	10552051
5	8675345
5	10144730

Table 3: "select AlbumId, AVG(Bytes) from track group by AlbumId"

AlbumId	AVG(Bytes)
1	6950574
2	5510424
3	4871098
4	9669293.5
5	9772926.25

```
dbGetQuery(chinook_db, "SELECT AlbumId, AVG(Bytes)
                        FROM Track
                        GROUP BY AlbumId
                        ORDER BY AVG(Bytes) DESC
                        LIMIT 10")
```

##	AlbumId	AVG(Bytes)
## 1	253	536359244
## 2	229	535292434
## 3	227	529469291
## 4	231	514373372
## 5	228	512231374
## 6	254	492670102
## 7	226	490750393
## 8	261	453454450
## 9	251	306109250
## 10	249	268393262

(Note: the order of commands here matters; try switching the order of GROUP BY and ORDER BY, you'll get an error)

We can use AS in the first line of SELECT to rename computed columns

```
dbGetQuery(chinook_db,  
            "SELECT AlbumId, AVG(Bytes) AS AvgBytes  
            FROM Track  
            GROUP BY AlbumId  
            ORDER BY AVG(Bytes) DESC  
            LIMIT 10")
```

```
##      AlbumId  AvgBytes  
## 1         253 536359244  
## 2         229 535292434  
## 3         227 529469291  
## 4         231 514373372  
## 5         228 512231374  
## 6         254 492670102  
## 7         226 490750393  
## 8         261 453454450  
## 9         251 306109250  
## 10        249 268393262
```


count, distinct, case insensitive

```
# case insensitive
```

```
dbGetQuery(chinook_db, "Select count(DISTINCT(ALBUMID)) FROM track")
```

```
## count(DISTINCT(ALBUMID))
```

```
## 1 347
```

"select AlbumId, Bytes from track"	
AlbumId	Bytes
16	12440200
16	8646737
35	5076048
50	39267613
50	29846063
67	9518842
67	8108507
67	11305791
67	9724150
67	8523388
67	7985133
35	21942829
35	13065612
50	27775442
67	6988128
67	8092463
35	17900787
35	9378873
50	24887209
67	8309725
67	10522352
67	10400020

"select AlbumId, Bytes from track Where AlbumId = 50"	
AlbumId	Bytes
50	39267613
50	29846063
50	27775442
50	24887209

"select AlbumId, AVG(Bytes) from track Where AlbumId = 50"	
AlbumId	AVG(Bytes)
50	30444081.75

Diagram illustrating data aggregation:

- Left table: "select AlbumId, Bytes from track" (Full track data)
- Middle table: "select AlbumId, Bytes from track Where AlbumId = 50" (Filtered data for AlbumId 50)
- Right table: "select AlbumId, AVG(Bytes) from track Where AlbumId = 50" (Aggregated data showing the average Bytes for AlbumId 50)
- Arrows indicate the flow of data from the full table to the filtered table, and then to the aggregated table.
- A label "Group average" points to the AVG(Bytes) column in the right table.

WHERE

```
dbGetQuery(chinook_db,  
            "SELECT  AlbumId, Avg(Bytes)  
              FROM Track  
              WHERE AlbumId = 50")
```

```
##   AlbumId Avg(Bytes)  
## 1      50   30444082
```

We can use the `WHERE` option in `SELECT` to specify a subset of the rows to use (*pre-aggregation/pre-calculation*)

```
dbGetQuery(chinook_db,  
            "SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes  
              FROM Track  
              WHERE AlbumId <= 160  
              GROUP BY AlbumId  
              ORDER BY AvgBytes DESC  
              LIMIT 10")
```

##	AlbumId	MediaTypeId	AvgBytes
## 1	50	1	30444082
## 2	138	1	24822832
## 3	137	1	19120969
## 4	43	1	16221538
## 5	97	1	16089011
## 6	114	1	15975057
## 7	109	1	15934275
## 8	113	1	15521017
## 9	127	1	15194926
## 10	98	1	14851676

Note we used the alias AvgBytes for AVG(BYTES) in the ORDER BY.

HAVING

We can use the HAVING option in SELECT to specify a subset of the rows to display (*post-aggregation/post-calculation*)

```
dbGetQuery(chinook_db,  
  "SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes  
  FROM Track  
  WHERE AlbumId >= 160  
  GROUP BY AlbumId  
  HAVING AvgBytes >= 25000000  
  ORDER BY AVG(Bytes) DESC  
  LIMIT 10")
```

##	AlbumId	MediaTypeId	AvgBytes
## 1	253	3	536359244
## 2	229	3	535292434
## 3	227	3	529469291
## 4	231	3	514373372
## 5	228	3	512231374
## 6	254	3	492670102
## 7	226	3	490750393
## 8	261	3	453454450
## 9	251	3	306109250
## 10	249	3	268393262

What happened?

```
dbGetQuery(chinook_db,
  paste("SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes",
        "FROM Track",
        "WHERE AlbumId >= 160"
  ))
```

```
##   AlbumId MediaTypeId AvgBytes
## 1      160           1 65784686
```

```
dbGetQuery(chinook_db,
  paste("SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes",
        "FROM Track",
        "WHERE AlbumId >= 160",
        "GROUP BY AlbumId"
  ))
```

```
##   AlbumId MediaTypeId AvgBytes
## 1      160           1 5948116
## 2      161           1 7287240
## 3      162           1 8528704
## 4      163           1 6228242
## 5      164           1 6937701
## 6      165           1 7830875
## 7      166           1 7153672
## 8      167           1 7127553
## 9      168           1 6939521
## 10     169           1 7485054
## 11     170           2 4601224
## 12     171           2 5390211
## 13     172           2 5179599
## 14     173           2 5832677
```

What happened?

```
dbGetQuery(chinook_db,  
            "SELECT AlbumId, MediaTypeId,AVG(Bytes) as AvgBytes  
            FROM Track  
            WHERE AlbumId >= 160  
            GROUP BY AlbumId  
            HAVING AvgBytes >= 25000000  
            ORDER BY AVG(Bytes) DESC  
            LIMIT 10")
```

##	AlbumId	MediaTypeId	AvgBytes
## 1	253	3	536359244
## 2	229	3	535292434
## 3	227	3	529469291
## 4	231	3	514373372
## 5	228	3	512231374
## 6	254	3	492670102
## 7	226	3	490750393
## 8	261	3	453454450
## 9	251	3	306109250
## 10	249	3	268393262

Pattern Matching on String Operators:

- **LIKE operator:** Allows wildcards to be used
 - `%` : matches any sequence of zero or more characters
 - `_` : matches any single character
- the `p%` pattern will match any strings that begin with `p`
 - e.g.: `pstat`, `pineapple`, `pop`
- the `%al` pattern matches any string that ends with `al`
 - e.g.: `pal`, `bridal`, `opal`
- the `%ul%` pattern matches any string that contains `ul`
 - e.g.: `ultimate` and `forceful`
- the pattern `r_n` will match the strings `run`, `ran`, `ron`
- the pattern `__rd` matches the strings `yard`, `ward`, `herd`, etc.
 - then what is the difference between `__rd` and `%rd` ?
 - `%rd` would also match longer words ending in `rd`, like `bernard`, `heard`, etc

Like operator

```
dbGetQuery(chinook_db,  
            "SELECT Name, Albumid, composer  
              FROM track  
              WHERE composer LIKE '%Smith%'  
              ORDER BY albumid  
              LIMIT 3")
```

```
##                Name AlbumId
```

```
## 1    Restless and Wild      3
```

```
## 2 Princess of the Dawn      3
```

```
## 3      Killing Floor     19
```

```
##
```

```
## 1 F. Baltes, R.A. Smith-Diesel, S. Kaufman, U. Dirkschneider &
```

```
## 2                                     Deaffy & R.A. S
```

```
## 3                                     A
```


IN operator

```
dbGetQuery(chinook_db,  
            "SELECT name, albumid, mediatypeid  
              FROM track  
              WHERE mediatypeid IN (2, 3)  
              LIMIT 5")
```

##	Name	AlbumId	MediaTypeId
## 1	Balls to the Wall	2	2
## 2	Fast As a Shark	3	2
## 3	Restless and Wild	3	2
## 4	Princess of the Dawn	3	2
## 5	Welcome to the Jungle	90	2

LIMIT offset

to get 10 rows starting from the 11th row in the track

```
dbGetQuery(chinook_db,  
            "SELECT trackId, name  
              FROM track  
              LIMIT 10 OFFSET 10")
```

##	TrackId	Name
## 1	11	C.O.D.
## 2	12	Breaking The Rules
## 3	13	Night Of The Long Knives
## 4	14	Spellbound
## 5	15	Go Down
## 6	16	Dog Eat Dog
## 7	17	Let There Be Rock
## 8	18	Bad Boy Boogie
## 9	19	Problem Child
## 10	20	Overdose

Disconnecting from the database

After the end of a session, it is good practice to explicitly close your connection.

```
dbDisconnect(chinook_db)

# Try selecting data
dbGetQuery(chinook_db,
           "select CustomerId, FirstName, LastName from Customer")

## Error: Invalid or closed connection
```

Does this remove the database connection `chinook_db` in the R session?

- All parts of a SQL query

SELECT columns or computations

FROM table

WHERE condition

GROUP BY columns

HAVING condition

ORDER BY column [ASC | DESC]

LIMIT offset, count;

- Database tools for R
 - the R packages RSQLite, DBI
 - the database Chinook_Sqlite.sqlite